

Chapter 06: Linear Model Selection and Regularization

Solutions to Exercises

February 03, 2023

CONCEPTUAL

EXERCISE 1:

Part a)

Best subset will have the smallest train RSS because the models will optimize on the training RSS and best subset will try every model that forward and backward selection will try.

Part b)

The best test RSS model could be any of the three. Best subset could easily overfit if the data has large p predictors relative to n observations. Forward and backward selection might not converge on the same model but try the same number of models and hard to say which selection process would be better.

Part c)

- i. TRUE
 - ii. TRUE
 - iii. FALSE
 - iv. FALSE
 - v. FALSE
-

EXERCISE 2:

Part a)

- iii. is TRUE - lasso puts a budget constraint on least squares (less flexible)

Part b)

- iii. is TRUE - ridge also puts a budget constraint on least squares (less flexible)

Part c)

- ii. is TRUE - a non-linear model would be more flexible and have higher variance, less bias
-

EXERCISE 3:

Part a)

- iv. is TRUE - as s is increased, there is less and less constraint on the model and it should always have a better training error (if s is increased to s' , then the best model using a budget of s would be included when using a budget of s')

Part b)

- ii. is TRUE - test error will improve (decrease) to a point and then will worsen (increase) as constraints loosen and model overfits

Part c)

- iii. is TRUE - variance always increases with fewer constraints

Part d)

- iv. is TRUE - bias always decreases with more model flexibility

Part e)

- v. is TRUE - the irreducible error is a constant value, not related to model selection
-

EXERCISE 4:

This problem is similar to Exercise 3, but for ridge instead of lasso and using λ instead of s . For each question part, ridge and lasso should be the same directionally except that increasing λ puts a heavier penalty in the equation, equivalent to reducing the budget s , so the answers to Exercise 4 should be flipped (horizontally) from answers in Exercise 3.

Part a)

- iii. is TRUE - training error increases steadily

Part b)

- ii. is TRUE - test error will decrease initially and then increase

Part c)

- iv. is TRUE - variance always decrease with more constraints

Part d)

iii. is TRUE - bias always increase with less model flexibility

Part e)

v. is TRUE - the irreducible error is a constant value, not related to model selection

EXERCISE 5:

Part a)

Ridge: minimize $(y_1 - \hat{\beta}_1 x_{11} - \hat{\beta}_2 x_{12})^2 + (y_2 - \hat{\beta}_1 x_{21} - \hat{\beta}_2 x_{22})^2 + \lambda(\hat{\beta}_1^2 + \hat{\beta}_2^2)$

Part b)

Step 1: Expanding the equation from Part a:

$$\begin{aligned} & (y_1 - \hat{\beta}_1 x_{11} - \hat{\beta}_2 x_{12})^2 + (y_2 - \hat{\beta}_1 x_{21} - \hat{\beta}_2 x_{22})^2 + \lambda(\hat{\beta}_1^2 + \hat{\beta}_2^2) \\ &= (y_1^2 + \hat{\beta}_1^2 x_{11}^2 + \hat{\beta}_2^2 x_{12}^2 - 2\hat{\beta}_1 x_{11} y_1 - 2\hat{\beta}_2 x_{12} y_1 + 2\hat{\beta}_1 \hat{\beta}_2 x_{11} x_{12}) \\ &+ (y_2^2 + \hat{\beta}_1^2 x_{21}^2 + \hat{\beta}_2^2 x_{22}^2 - 2\hat{\beta}_1 x_{21} y_2 - 2\hat{\beta}_2 x_{22} y_2 + 2\hat{\beta}_1 \hat{\beta}_2 x_{21} x_{22}) \\ &+ \lambda\hat{\beta}_1^2 + \lambda\hat{\beta}_2^2 \end{aligned}$$

Step 2: Taking the partial derivative to $\hat{\beta}_1$ and setting equation to 0 to minimize:

$$\frac{\partial}{\partial \hat{\beta}_1} : (2\hat{\beta}_1 x_{11}^2 - 2x_{11} y_1 + 2\hat{\beta}_2 x_{11} x_{12}) + (2\hat{\beta}_1 x_{21}^2 - 2x_{21} y_2 + 2\hat{\beta}_2 x_{21} x_{22}) + 2\lambda\hat{\beta}_1 = 0$$

Step 3: Setting $x_{11} = x_{12} = x_1$ and $x_{21} = x_{22} = x_2$ and dividing both sides of the equation by 2:

$$(\hat{\beta}_1 x_1^2 - x_1 y_1 + \hat{\beta}_2 x_1^2) + (\hat{\beta}_1 x_2^2 - x_2 y_2 + \hat{\beta}_2 x_2^2) + \lambda\hat{\beta}_1 = 0$$

$$\hat{\beta}_1(x_1^2 + x_2^2) + \hat{\beta}_2(x_1^2 + x_2^2) + \lambda\hat{\beta}_1 = x_1 y_1 + x_2 y_2$$

Step 4: Add $2\hat{\beta}_1 x_1 x_2$ and $2\hat{\beta}_2 x_1 x_2$ to both sides of the equation:

$$\begin{aligned} & \hat{\beta}_1(x_1^2 + x_2^2 + 2x_1 x_2) + \hat{\beta}_2(x_1^2 + x_2^2 + 2x_1 x_2) + \lambda\hat{\beta}_1 = x_1 y_1 + x_2 y_2 + 2\hat{\beta}_1 x_1 x_2 + 2\hat{\beta}_2 x_1 x_2 \\ & \hat{\beta}_1(x_1 + x_2)^2 + \hat{\beta}_2(x_1 + x_2)^2 + \lambda\hat{\beta}_1 = x_1 y_1 + x_2 y_2 + 2\hat{\beta}_1 x_1 x_2 + 2\hat{\beta}_2 x_1 x_2 \end{aligned}$$

Step 5: Because $x_1 + x_2 = 0$, we can eliminate the first two terms:

$$\lambda\hat{\beta}_1 = x_1 y_1 + x_2 y_2 + 2\hat{\beta}_1 x_1 x_2 + 2\hat{\beta}_2 x_1 x_2$$

Step 6: Similarly by taking the partial derivative to $\hat{\beta}_2$, we can get the equation:

$$\lambda\hat{\beta}_2 = x_1 y_1 + x_2 y_2 + 2\hat{\beta}_1 x_1 x_2 + 2\hat{\beta}_2 x_1 x_2$$

Step 7: The left side of the equations for both $\lambda\hat{\beta}_1$ and $\lambda\hat{\beta}_2$ are the same so we have:

$$\lambda \hat{\beta}_1 = \lambda \hat{\beta}_2$$

$$\hat{\beta}_1 = \hat{\beta}_2$$

Part c)

Lasso: minimize $(y_1 - \hat{\beta}_1 x_{11} - \hat{\beta}_2 x_{12})^2 + (y_2 - \hat{\beta}_1 x_{21} - \hat{\beta}_2 x_{22})^2 + \lambda(|\hat{\beta}_1| + |\hat{\beta}_2|)$

Part d)

Replacing the constraint term from Part b, the derivative term to β is:

$$\frac{\partial}{\partial \hat{\beta}}(\lambda|\beta|) : \lambda \frac{|\beta|}{\beta}$$

Following through the steps in Part b, we get:

$$\lambda \frac{|\beta_1|}{\beta_1} = \lambda \frac{|\beta_2|}{\beta_2}$$


So it seems that the lasso just requires that β_1 and β_2 are both positive or both negative (ignoring possibility of 0...)

EXERCISE 6:

Part a)

```
betas <- seq(-10,10,0.1)
eq.ridge <- function(beta, y=7, lambda=10) (y-beta)^2 + lambda*beta^2
plot(betas, eq.ridge(betas), xlab="beta", main="Ridge Regression Optimization", pch=1)
points(5/(1+10), eq.ridge(7/(1+10)), pch=16, col="red", cex=2)
```

Ridge Regression Optimization




For $y = 7$ and $\lambda = 10$, $\hat{\beta} = \frac{7}{1+10}$ minimizes the ridge regression equation

Part b)

```
betas <- seq(-10,10,0.1)
eq.lasso <- function(beta, y=7, lambda=10) (y-beta)^2 + lambda*abs(beta)
plot(betas, eq.lasso(betas), xlab="beta", main="Lasso Regression Optimization", pch=1)
points(7-10/2, eq.lasso(7-10/2), pch=16, col="red", cex=2)
```

Lasso Regression Optimization



For $y = 7$ and $\lambda = 10$, $\hat{\beta} = 7 - \frac{10}{2}$ minimizes the ridge regression equation

EXERCISE 7:

Part a)

[... will come back to this. maybe.]

Part b)

[... will come back to this. maybe.]

Part c)

[... will come back to this. maybe.]

Part d)

[... will come back to this. maybe.]

Part e)

[... will come back to this. maybe.]

APPLIED


EXERCISE 8:

Part a)

```
set.seed(1)
X <- rnorm(100)
eps <- rnorm(100)
```

Part b)

```
Y <- 5 + 10*X - 3*X^2 + 2*X^3 + eps
plot(X,Y)
```



Part c)

```
require(leaps)
regfit.full <- regsubsets(Y~poly(X,10,raw=T), data=data.frame(Y,X), nvmax=10)
(reg.summary <- summary(regfit.full))
```

```
## Subset selection object
```

```
## Call: regsubsets.formula(Y ~ poly(X, 10, raw = T), data = data.frame(Y,
```

```

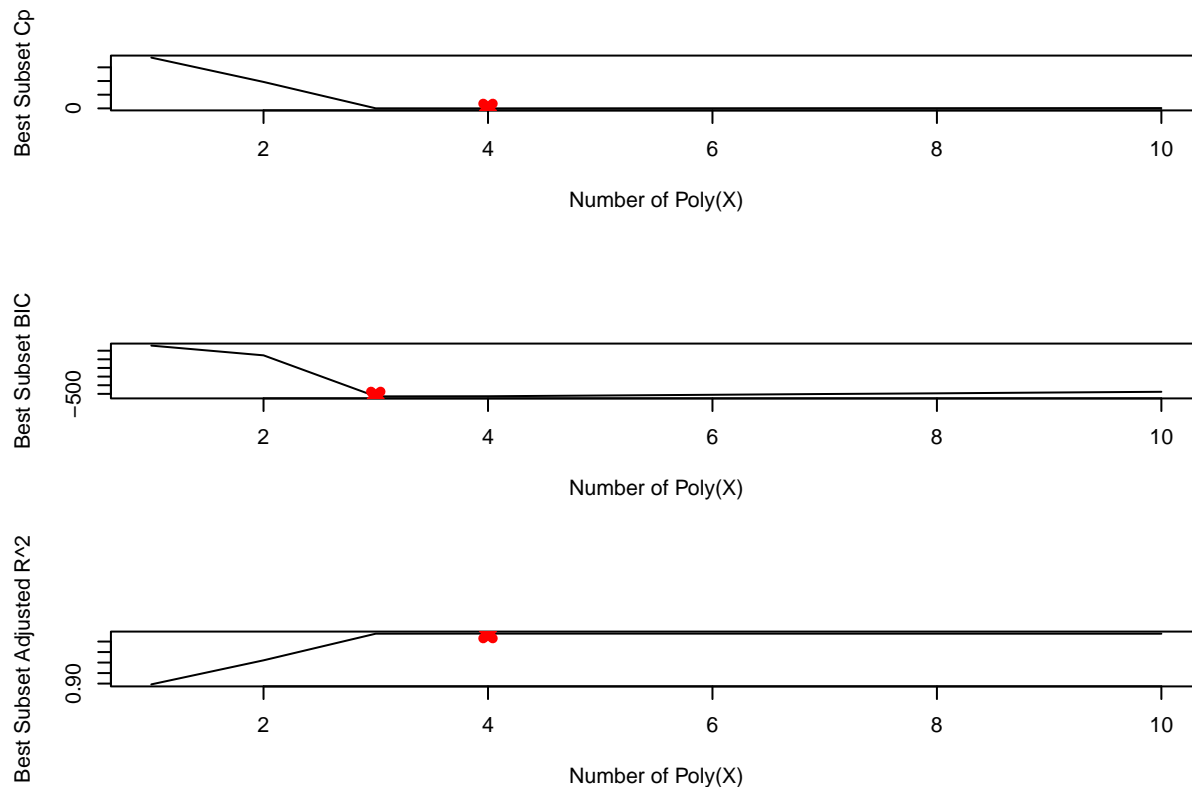
##      X), nvmax = 10)
## 10 Variables (and intercept)
##              Forced in Forced out
## poly(X, 10, raw = T)1      FALSE      FALSE
## poly(X, 10, raw = T)2      FALSE      FALSE
## poly(X, 10, raw = T)3      FALSE      FALSE
## poly(X, 10, raw = T)4      FALSE      FALSE
## poly(X, 10, raw = T)5      FALSE      FALSE
## poly(X, 10, raw = T)6      FALSE      FALSE
## poly(X, 10, raw = T)7      FALSE      FALSE
## poly(X, 10, raw = T)8      FALSE      FALSE
## poly(X, 10, raw = T)9      FALSE      FALSE
## poly(X, 10, raw = T)10     FALSE      FALSE
## 1 subsets of each size up to 10
## Selection Algorithm: exhaustive
##      poly(X, 10, raw = T)1 poly(X, 10, raw = T)2 poly(X, 10, raw = T)3
## 1 ( 1 ) "*"          " "          " "
## 2 ( 1 ) "*"          "*"          " "
## 3 ( 1 ) "*"          "*"          "*"
## 4 ( 1 ) "*"          "*"          "*"
## 5 ( 1 ) "*"          "*"          "*"
## 6 ( 1 ) "*"          "*"          "*"
## 7 ( 1 ) "*"          "*"          "*"
## 8 ( 1 ) "*"          "*"          "*"
## 9 ( 1 ) "*"          "*"          "*"
## 10 ( 1 ) "*"         "*"          "*"
##      poly(X, 10, raw = T)4 poly(X, 10, raw = T)5 poly(X, 10, raw = T)6
## 1 ( 1 ) " "          " "          " "
## 2 ( 1 ) " "          " "          " "
## 3 ( 1 ) " "          " "          " "
## 4 ( 1 ) " "          "*"          " "
## 5 ( 1 ) " "          "*"          "*"
## 6 ( 1 ) " "          " "          " "
## 7 ( 1 ) " "          "*"          "*"
## 8 ( 1 ) "*"          " "          "*"
## 9 ( 1 ) "*"          "*"          "*"
## 10 ( 1 ) "*"         "*"          "*"
##      poly(X, 10, raw = T)7 poly(X, 10, raw = T)8 poly(X, 10, raw = T)9
## 1 ( 1 ) " "          " "          " "
## 2 ( 1 ) " "          " "          " "
## 3 ( 1 ) " "          " "          " "
## 4 ( 1 ) " "          " "          " "
## 5 ( 1 ) " "          " "          " "
## 6 ( 1 ) "*"          "*"          "*"
## 7 ( 1 ) " "          "*"          " "
## 8 ( 1 ) " "          "*"          "*"
## 9 ( 1 ) " "          "*"          "*"
## 10 ( 1 ) "*"         "*"          "*"
##      poly(X, 10, raw = T)10
## 1 ( 1 ) " "
## 2 ( 1 ) " "
## 3 ( 1 ) " "
## 4 ( 1 ) " "
## 5 ( 1 ) " "

```



```
## 6 ( 1 ) " "
## 7 ( 1 ) "*"
## 8 ( 1 ) "*"
## 9 ( 1 ) "*"
## 10 ( 1 ) "*"
```

```
par(mfrow=c(3,1))
min.cp <- which.min(reg.summary$cp)
plot(reg.summary$cp, xlab="Number of Poly(X)", ylab="Best Subset Cp", type="l")
points(min.cp, reg.summary$cp[min.cp], col="red", pch=4, lwd=5)
min.bic <- which.min(reg.summary$bic)
plot(reg.summary$bic, xlab="Number of Poly(X)", ylab="Best Subset BIC", type="l")
points(min.bic, reg.summary$bic[min.bic], col="red", pch=4, lwd=5)
min.adj2 <- which.max(reg.summary$adjr2)
plot(reg.summary$adjr2, xlab="Number of Poly(X)", ylab="Best Subset Adjusted R^2", type="l")
points(min.adj2, reg.summary$adjr2[min.adj2], col="red", pch=4, lwd=5)
```



```
coef(regfit.full, min.cp)
```

```
##          (Intercept) poly(X, 10, raw = T)1 poly(X, 10, raw = T)2
##          5.07200775      10.38745596      -3.15424359
## poly(X, 10, raw = T)3 poly(X, 10, raw = T)5
##          1.55797426      0.08072292
```

```
coef(regfit.full, min.bic)
```

```
##          (Intercept) poly(X, 10, raw = T)1 poly(X, 10, raw = T)2
##          5.061507          9.975280          -3.123791
## poly(X, 10, raw = T)3
##          2.017639
```

```
coef(regfit.full, min.adjR2)
```

```
##          (Intercept) poly(X, 10, raw = T)1 poly(X, 10, raw = T)2
##          5.07200775          10.38745596          -3.15424359
## poly(X, 10, raw = T)3 poly(X, 10, raw = T)5
##          1.55797426          0.08072292
```

Part d)

```
# forward selection
```

```
regfit.fwd <- regsubsets(Y~poly(X,10,raw=T), data=data.frame(Y,X), nvmax=10)
(fwd.summary <- summary(regfit.fwd))
```

```
## Subset selection object
## Call: regsubsets.formula(Y ~ poly(X, 10, raw = T), data = data.frame(Y,
##      X), nvmax = 10)
## 10 Variables (and intercept)
##              Forced in Forced out
## poly(X, 10, raw = T)1      FALSE      FALSE
## poly(X, 10, raw = T)2      FALSE      FALSE
## poly(X, 10, raw = T)3      FALSE      FALSE
## poly(X, 10, raw = T)4      FALSE      FALSE
## poly(X, 10, raw = T)5      FALSE      FALSE
## poly(X, 10, raw = T)6      FALSE      FALSE
## poly(X, 10, raw = T)7      FALSE      FALSE
## poly(X, 10, raw = T)8      FALSE      FALSE
## poly(X, 10, raw = T)9      FALSE      FALSE
## poly(X, 10, raw = T)10     FALSE      FALSE
## 1 subsets of each size up to 10
## Selection Algorithm: exhaustive
##          poly(X, 10, raw = T)1 poly(X, 10, raw = T)2 poly(X, 10, raw = T)3
## 1  ( 1 )  "*"                " "                " "
## 2  ( 1 )  "*"                "*"                " "
## 3  ( 1 )  "*"                "*"                "*"
## 4  ( 1 )  "*"                "*"                "*"
## 5  ( 1 )  "*"                "*"                "*"
## 6  ( 1 )  "*"                "*"                "*"
## 7  ( 1 )  "*"                "*"                "*"
## 8  ( 1 )  "*"                "*"                "*"
## 9  ( 1 )  "*"                "*"                "*"
## 10 ( 1 )  "*"                "*"                "*"
##          poly(X, 10, raw = T)4 poly(X, 10, raw = T)5 poly(X, 10, raw = T)6
## 1  ( 1 )  " "                " "                " "
## 2  ( 1 )  " "                " "                " "
```

```
## 3 ( 1 ) " " " "
## 4 ( 1 ) " " "*" " "
## 5 ( 1 ) " " "*" "*"
## 6 ( 1 ) " " " " " "
## 7 ( 1 ) " " "*" "*"
## 8 ( 1 ) "*" " " "*"
## 9 ( 1 ) "*" "*" "*"
## 10 ( 1 ) "*" "*" "*"
##      poly(X, 10, raw = T)7 poly(X, 10, raw = T)8 poly(X, 10, raw = T)9
## 1 ( 1 ) " " " " " "
## 2 ( 1 ) " " " " " "
## 3 ( 1 ) " " " " " "
## 4 ( 1 ) " " " " " "
## 5 ( 1 ) " " " " " "
## 6 ( 1 ) "*" "*" "*"
## 7 ( 1 ) " " "*" " "
## 8 ( 1 ) " " "*" "*"
## 9 ( 1 ) " " "*" "*"
## 10 ( 1 ) "*" "*" "*"
##      poly(X, 10, raw = T)10
## 1 ( 1 ) " "
## 2 ( 1 ) " "
## 3 ( 1 ) " "
## 4 ( 1 ) " "
## 5 ( 1 ) " "
## 6 ( 1 ) " "
## 7 ( 1 ) "*"
## 8 ( 1 ) "*"
## 9 ( 1 ) "*"
## 10 ( 1 ) "*"

```

```
# backward selection
```

```
regfit.bwd <- regsubsets(Y~poly(X,10,raw=T), data=data.frame(Y,X), nvmax=10)
(bwd.summary <- summary(regfit.bwd))
```

```
## Subset selection object
## Call: regsubsets.formula(Y ~ poly(X, 10, raw = T), data = data.frame(Y,
##      X), nvmax = 10)
## 10 Variables (and intercept)
##      Forced in Forced out
## poly(X, 10, raw = T)1 FALSE FALSE
## poly(X, 10, raw = T)2 FALSE FALSE
## poly(X, 10, raw = T)3 FALSE FALSE
## poly(X, 10, raw = T)4 FALSE FALSE
## poly(X, 10, raw = T)5 FALSE FALSE
## poly(X, 10, raw = T)6 FALSE FALSE
## poly(X, 10, raw = T)7 FALSE FALSE
## poly(X, 10, raw = T)8 FALSE FALSE
## poly(X, 10, raw = T)9 FALSE FALSE
## poly(X, 10, raw = T)10 FALSE FALSE
## 1 subsets of each size up to 10
## Selection Algorithm: exhaustive
##      poly(X, 10, raw = T)1 poly(X, 10, raw = T)2 poly(X, 10, raw = T)3
## 1 ( 1 ) "*" " " " "
```

```

## 2 ( 1 ) "*" "*" " "
## 3 ( 1 ) "*" "*" "*"
## 4 ( 1 ) "*" "*" "*"
## 5 ( 1 ) "*" "*" "*"
## 6 ( 1 ) "*" "*" "*"
## 7 ( 1 ) "*" "*" "*"
## 8 ( 1 ) "*" "*" "*"
## 9 ( 1 ) "*" "*" "*"
## 10 ( 1 ) "*" "*" "*"
##      poly(X, 10, raw = T)4 poly(X, 10, raw = T)5 poly(X, 10, raw = T)6
## 1 ( 1 ) " " " " " "
## 2 ( 1 ) " " " " " "
## 3 ( 1 ) " " " " " "
## 4 ( 1 ) " " "*" " "
## 5 ( 1 ) " " "*" "*"
## 6 ( 1 ) " " " " " "
## 7 ( 1 ) " " "*" "*"
## 8 ( 1 ) "*" " " "*"
## 9 ( 1 ) "*" "*" "*"
## 10 ( 1 ) "*" "*" "*"
##      poly(X, 10, raw = T)7 poly(X, 10, raw = T)8 poly(X, 10, raw = T)9
## 1 ( 1 ) " " " " " "
## 2 ( 1 ) " " " " " "
## 3 ( 1 ) " " " " " "
## 4 ( 1 ) " " " " " "
## 5 ( 1 ) " " " " " "
## 6 ( 1 ) "*" "*" "*"
## 7 ( 1 ) " " "*" " "
## 8 ( 1 ) " " "*" "*"
## 9 ( 1 ) " " "*" "*"
## 10 ( 1 ) "*" "*" "*"
##      poly(X, 10, raw = T)10
## 1 ( 1 ) " "
## 2 ( 1 ) " "
## 3 ( 1 ) " "
## 4 ( 1 ) " "
## 5 ( 1 ) " "
## 6 ( 1 ) " "
## 7 ( 1 ) "*"
## 8 ( 1 ) "*"
## 9 ( 1 ) "*"
## 10 ( 1 ) "*"

```

```

par(mfrow=c(3,2))

min.cp <- which.min(fwd.summary$cp)
plot(fwd.summary$cp, xlab="Number of Poly(X)", ylab="Forward Selection Cp", type="l")
points(min.cp, fwd.summary$cp[min.cp], col="red", pch=4, lwd=5)

min.cp <- which.min(bwd.summary$cp)
plot(bwd.summary$cp, xlab="Number of Poly(X)", ylab="Backward Selection Cp", type="l")
points(min.cp, bwd.summary$cp[min.cp], col="red", pch=4, lwd=5)

min.bic <- which.min(fwd.summary$bic)

```

```

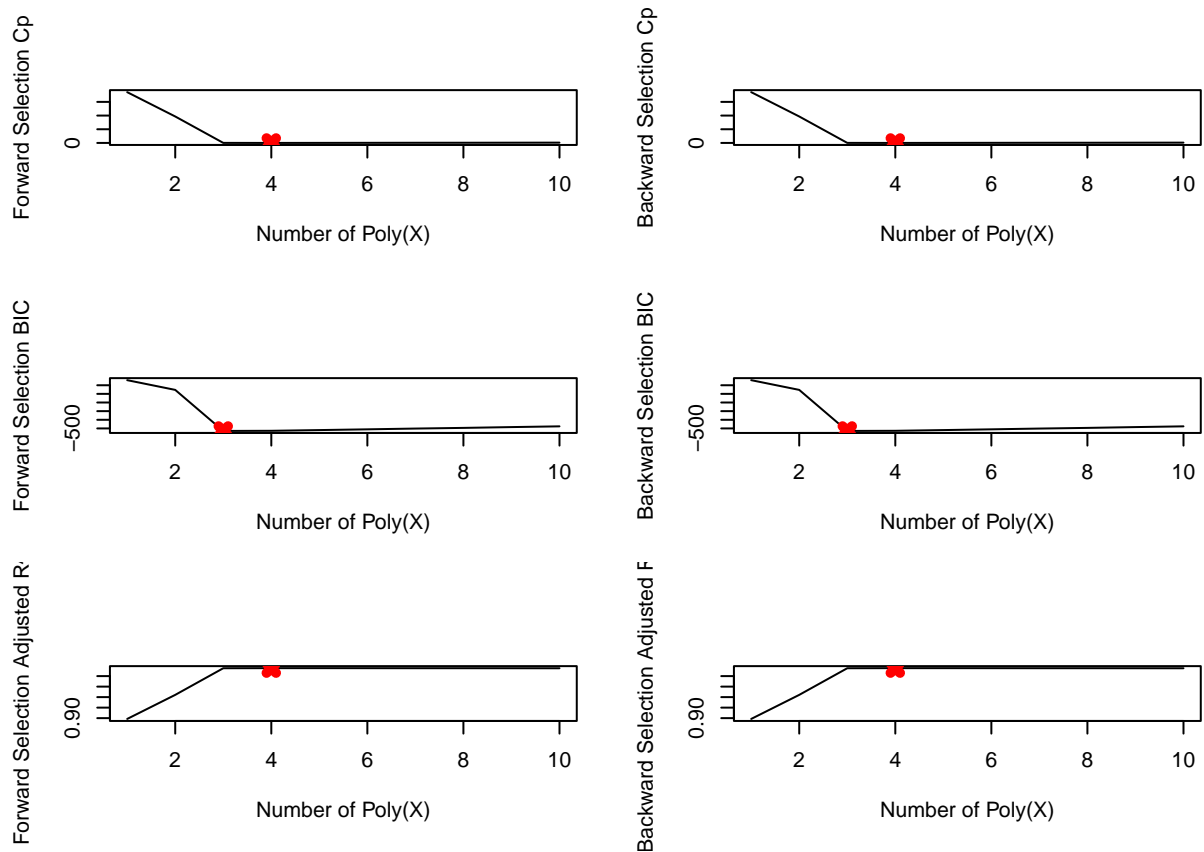
plot(fwd.summary$bic, xlab="Number of Poly(X)", ylab="Forward Selection BIC", type="l")
points(min.bic, fwd.summary$bic[min.bic], col="red", pch=4, lwd=5)

min.bic <- which.min(bwd.summary$bic)
plot(bwd.summary$bic, xlab="Number of Poly(X)", ylab="Backward Selection BIC", type="l")
points(min.bic, bwd.summary$bic[min.bic], col="red", pch=4, lwd=5)

min.adj2 <- which.max(fwd.summary$adj2)
plot(fwd.summary$adj2, xlab="Number of Poly(X)", ylab="Forward Selection Adjusted R^2", type="l")
points(min.adj2, fwd.summary$adj2[min.adj2], col="red", pch=4, lwd=5)

min.adj2 <- which.max(bwd.summary$adj2)
plot(bwd.summary$adj2, xlab="Number of Poly(X)", ylab="Backward Selection Adjusted R^2", type="l")
points(min.adj2, bwd.summary$adj2[min.adj2], col="red", pch=4, lwd=5)

```



```

# coefficients of selected models
coef(regfit.fwd, which.min(fwd.summary$cp))

```

```

##          (Intercept) poly(X, 10, raw = T)1 poly(X, 10, raw = T)2
##          5.07200775      10.38745596      -3.15424359
## poly(X, 10, raw = T)3 poly(X, 10, raw = T)5
##          1.55797426      0.08072292

```

```
coef(regfit.bwd, which.min(bwd.summary$cp))
```

```
##          (Intercept) poly(X, 10, raw = T)1 poly(X, 10, raw = T)2
##          5.07200775      10.38745596      -3.15424359
## poly(X, 10, raw = T)3 poly(X, 10, raw = T)5
##          1.55797426      0.08072292
```

```
coef(regfit.fwd, which.min(fwd.summary$bic))
```

```
##          (Intercept) poly(X, 10, raw = T)1 poly(X, 10, raw = T)2
##          5.061507      9.975280      -3.123791
## poly(X, 10, raw = T)3
##          2.017639
```

```
coef(regfit.bwd, which.min(bwd.summary$bic))
```

```
##          (Intercept) poly(X, 10, raw = T)1 poly(X, 10, raw = T)2
##          5.061507      9.975280      -3.123791
## poly(X, 10, raw = T)3
##          2.017639
```

```
coef(regfit.fwd, which.max(fwd.summary$adjr2))
```

```
##          (Intercept) poly(X, 10, raw = T)1 poly(X, 10, raw = T)2
##          5.07200775      10.38745596      -3.15424359
## poly(X, 10, raw = T)3 poly(X, 10, raw = T)5
##          1.55797426      0.08072292
```

```
coef(regfit.bwd, which.max(bwd.summary$adjr2))
```

```
##          (Intercept) poly(X, 10, raw = T)1 poly(X, 10, raw = T)2
##          5.07200775      10.38745596      -3.15424359
## poly(X, 10, raw = T)3 poly(X, 10, raw = T)5
##          1.55797426      0.08072292
```


Best subset, forward selection and backward selection all resulted in the same best models

Part e)

```
require(glmnet)
xmat <- model.matrix(Y~poly(X,10,raw=T))[,,-1]
lasso.mod <- cv.glmnet(xmat, Y, alpha=1)
(lambda <- lasso.mod$lambda.min)
```

```
## [1] 0.04835977
```

```
par(mfrow=c(1,1))
plot(lasso.mod)
```



```
predict(lasso.mod, s=lambda, type="coefficients")
```

```
## 11 x 1 sparse Matrix of class "dgCMatrix"
##                                s1
## (Intercept)                   5.032793082
## poly(X, 10, raw = T)1         10.211994502
## poly(X, 10, raw = T)2        -3.092418697
## poly(X, 10, raw = T)3         1.708546779
## poly(X, 10, raw = T)4         .
## poly(X, 10, raw = T)5         0.049655684
## poly(X, 10, raw = T)6         .
## poly(X, 10, raw = T)7         0.000459821
## poly(X, 10, raw = T)8         .
## poly(X, 10, raw = T)9         .
## poly(X, 10, raw = T)10        .
```

Lasso regression selects the correct predictors: X , X^2 and X^3

Part f)

```
Y2 <- 5 + 1.5*X^7 + eps

# best subset model selection
regfit.full <- regsubsets(Y2~poly(X,10,raw=T), data=data.frame(Y,X), nvmax=10)
par(mfrow=c(3,1))
(reg.summary <- summary(regfit.full))
```

```




## Subset selection object
## Call: regsubsets.formula(Y2 ~ poly(X, 10, raw = T), data = data.frame(Y,
##      X), nvmax = 10)
## 10 Variables (and intercept)
##              Forced in Forced out
## poly(X, 10, raw = T)1      FALSE      FALSE
## poly(X, 10, raw = T)2      FALSE      FALSE
## poly(X, 10, raw = T)3      FALSE      FALSE
## poly(X, 10, raw = T)4      FALSE      FALSE
## poly(X, 10, raw = T)5      FALSE      FALSE
## poly(X, 10, raw = T)6      FALSE      FALSE
## poly(X, 10, raw = T)7      FALSE      FALSE
## poly(X, 10, raw = T)8      FALSE      FALSE
## poly(X, 10, raw = T)9      FALSE      FALSE
## poly(X, 10, raw = T)10     FALSE      FALSE
## 1 subsets of each size up to 10
## Selection Algorithm: exhaustive
##      poly(X, 10, raw = T)1 poly(X, 10, raw = T)2 poly(X, 10, raw = T)3
## 1  ( 1 )  " "          " "          " "
## 2  ( 1 )  " "          "*"          " "
## 3  ( 1 )  " "          "*"          " "
## 4  ( 1 )  "*"          "*"          "*"
## 5  ( 1 )  "*"          "*"          "*"
## 6  ( 1 )  "*"          " "          "*"
## 7  ( 1 )  "*"          " "          "*"
## 8  ( 1 )  "*"          "*"          "*"
## 9  ( 1 )  "*"          "*"          "*"
## 10 ( 1 )  "*"          "*"          "*"
##      poly(X, 10, raw = T)4 poly(X, 10, raw = T)5 poly(X, 10, raw = T)6
## 1  ( 1 )  " "          " "          " "
## 2  ( 1 )  " "          " "          " "
## 3  ( 1 )  " "          "*"          " "
## 4  ( 1 )  " "          " "          " "
## 5  ( 1 )  "*"          " "          " "
## 6  ( 1 )  " "          " "          "*"
## 7  ( 1 )  " "          "*"          "*"
## 8  ( 1 )  "*"          " "          "*"
## 9  ( 1 )  "*"          " "          "*"
## 10 ( 1 )  "*"          "*"          "*"
##      poly(X, 10, raw = T)7 poly(X, 10, raw = T)8 poly(X, 10, raw = T)9
## 1  ( 1 )  "*"          " "          " "
## 2  ( 1 )  "*"          " "          " "
## 3  ( 1 )  "*"          " "          " "
## 4  ( 1 )  "*"          " "          " "
## 5  ( 1 )  "*"          " "          " "
## 6  ( 1 )  "*"          "*"          " "
## 7  ( 1 )  "*"          "*"          " "
## 8  ( 1 )  "*"          "*"          " "
## 9  ( 1 )  "*"          "*"          "*"
## 10 ( 1 )  "*"          "*"          "*"
##      poly(X, 10, raw = T)10
## 1  ( 1 )  " "
## 2  ( 1 )  " "
## 3  ( 1 )  " "

```



```
## 4 ( 1 ) " "
## 5 ( 1 ) " "
## 6 ( 1 ) "*"
## 7 ( 1 ) "*"
## 8 ( 1 ) "*"
## 9 ( 1 ) "*"
## 10 ( 1 ) "*"
```

```
min.cp <- which.min(reg.summary$cp)
plot(reg.summary$cp, xlab="Number of Poly(X)", ylab="Best Subset Cp", type="l")
points(min.cp, reg.summary$cp[min.cp], col="red", pch=4, lwd=5)
min.bic <- which.min(reg.summary$bic)
plot(reg.summary$bic, xlab="Number of Poly(X)", ylab="Best Subset BIC", type="l")
points(min.bic, reg.summary$bic[min.bic], col="red", pch=4, lwd=5)
min.adj2 <- which.max(reg.summary$adj2)
plot(reg.summary$adj2, xlab="Number of Poly(X)", ylab="Best Subset Adjusted R^2", type="l")
points(min.adj2, reg.summary$adj2[min.adj2], col="red", pch=4, lwd=5)
```



```
coef(regfit.full, min.cp)
```

```
## (Intercept) poly(X, 10, raw = T)2 poly(X, 10, raw = T)7
## 5.0704904 -0.1417084 1.5015552
```

```
coef(regfit.full, min.bic)
```

```
##          (Intercept) poly(X, 10, raw = T)7
##          4.95894          1.50077
```

```
coef(regfit.full, min.adjr2)
```


```
##          (Intercept) poly(X, 10, raw = T)1 poly(X, 10, raw = T)2
##          5.0762524          0.2914016          -0.1617671
## poly(X, 10, raw = T)3 poly(X, 10, raw = T)7
##          -0.2526527          1.5091338
```

```
# lasso regression
```

```
xmat <- model.matrix(Y2~poly(X,10,raw=T))[, -1]
lasso.mod <- cv.glmnet(xmat, Y2, alpha=1)
(lambda <- lasso.mod$lambda.min)
```

```
## [1] 2.651535
```

```
par(mfrow=c(1,1))
plot(lasso.mod)
```



```
predict(lasso.mod, s=lambda, type="coefficients")
```

```
## 11 x 1 sparse Matrix of class "dgCMatrix"
##                               s1
## (Intercept)                5.143574
## poly(X, 10, raw = T)1      .
## poly(X, 10, raw = T)2      .
## poly(X, 10, raw = T)3      .
## poly(X, 10, raw = T)4      .
## poly(X, 10, raw = T)5      .
## poly(X, 10, raw = T)6      .
## poly(X, 10, raw = T)7      1.457022
## poly(X, 10, raw = T)8      .
## poly(X, 10, raw = T)9      .
## poly(X, 10, raw = T)10     .
```

Lasso selects the correct model but best subset diagnostics indicate using 1 to 4 predictors

EXERCISE 9:

Part a)

```
require(ISLR2)
data(College)
set.seed(1)
trainid <- sample(1:nrow(College), nrow(College)/2)
train <- College[trainid,]
test <- College[-trainid,]
str(College)
```

```
## 'data.frame': 777 obs. of 18 variables:
## $ Private : Factor w/ 2 levels "No","Yes": 2 2 2 2 2 2 2 2 2 2 ...
## $ Apps : num 1660 2186 1428 417 193 ...
## $ Accept : num 1232 1924 1097 349 146 ...
## $ Enroll : num 721 512 336 137 55 158 103 489 227 172 ...
## $ Top10perc : num 23 16 22 60 16 38 17 37 30 21 ...
## $ Top25perc : num 52 29 50 89 44 62 45 68 63 44 ...
## $ F.Undergrad: num 2885 2683 1036 510 249 ...
## $ P.Undergrad: num 537 1227 99 63 869 ...
## $ Outstate : num 7440 12280 11250 12960 7560 ...
## $ Room.Board : num 3300 6450 3750 5450 4120 ...
## $ Books : num 450 750 400 450 800 500 500 450 300 660 ...
## $ Personal : num 2200 1500 1165 875 1500 ...
## $ PhD : num 70 29 53 92 76 67 90 89 79 40 ...
## $ Terminal : num 78 30 66 97 72 73 93 100 84 41 ...
## $ S.F.Ratio : num 18.1 12.2 12.9 7.7 11.9 9.4 11.5 13.7 11.3 11.5 ...
## $ perc.alumni: num 12 16 30 37 2 11 26 37 23 15 ...
## $ Expend : num 7041 10527 8735 19016 10922 ...
## $ Grad.Rate : num 60 56 54 59 15 55 63 73 80 52 ...
```

Part b)

```
fit.lm <- lm(Apps~., data=train)
pred.lm <- predict(fit.lm, test)
(err.lm <- mean((test$Apps - pred.lm)^2)) # test error
```

```
## [1] 1135758
```

Part c)

```
require(glmnet)
xmat.train <- model.matrix(Apps~., data=train)[,-1]
xmat.test <- model.matrix(Apps~., data=test)[,-1]
fit.ridge <- cv.glmnet(xmat.train, train$Apps, alpha=0)
(lambda <- fit.ridge$lambda.min) # optimal lambda
```

```
## [1] 405.8404
```

```
pred.ridge <- predict(fit.ridge, s=lambda, newx=xmat.test)
(err.ridge <- mean((test$Apps - pred.ridge)^2)) # test error
```

```
## [1] 976261.5
```

Part d)

```
require(glmnet)
xmat.train <- model.matrix(Apps~., data=train)[,-1]
xmat.test <- model.matrix(Apps~., data=test)[,-1]
fit.lasso <- cv.glmnet(xmat.train, train$Apps, alpha=1)
(lambda <- fit.lasso$lambda.min) # optimal lambda
```

```
## [1] 1.97344
```

```
pred.lasso <- predict(fit.lasso, s=lambda, newx=xmat.test)
(err.lasso <- mean((test$Apps - pred.lasso)^2)) # test error
```

```
## [1] 1115901
```

```
coef.lasso <- predict(fit.lasso, type="coefficients", s=lambda)[1:ncol(College),]
coef.lasso[coef.lasso != 0]
```


```
##      (Intercept)      PrivateYes      Accept      Enroll      Top10perc
## -7.688896e+02 -3.127034e+02  1.762718e+00 -1.318195e+00  6.482356e+01
##      Top25perc      F.Undergrad      P.Undergrad      Outstate      Room.Board
## -2.081406e+01  7.119149e-02  1.246161e-02 -1.049091e-01  2.088305e-01
##      Books      Personal      PhD      Terminal      S.F.Ratio
##  2.926466e-01  3.955068e-03 -1.455463e+01  5.395858e+00  2.171398e+01
##      perc.alumni      Expend      Grad.Rate
##  5.088260e-01  4.824455e-02  7.036148e+00
```

```
length(coef.lasso[coef.lasso != 0])
```

```
## [1] 18
```

Part e)

```
require(pls)
set.seed(1)
fit.pcr <- pcr(Apps~., data=train, scale=TRUE, validation="CV")
validationplot(fit.pcr, val.type="MSEP")
```



```
summary(fit.pcr)
```

```
## Data:      X dimension: 388 17
## Y dimension: 388 1
## Fit method: svdpc
## Number of components considered: 17
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV           4288    4006    2373    2372    2069    1961    1919
## adjCV         4288    4007    2368    2369    1999    1948    1911
##      7 comps  8 comps  9 comps 10 comps 11 comps 12 comps 13 comps
```

```
## CV      1919      1921      1876      1832      1832      1836      1837
## adjCV    1912      1915      1868      1821      1823      1827      1827
##      14 comps  15 comps  16 comps  17 comps
## CV      1853      1759      1341      1270
## adjCV    1850      1733      1326      1257
##
## TRAINING: % variance explained
##      1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps  8 comps
## X      32.20   57.78   65.31   70.99   76.37   81.27   84.8    87.85
## Apps   13.44   70.93   71.07   79.87   81.15   82.25   82.3    82.33
##      9 comps  10 comps 11 comps 12 comps 13 comps 14 comps 15 comps
## X      90.62   92.91   94.98   96.74   97.79   98.72   99.42
## Apps   83.38   84.76   84.80   84.84   85.11   85.14   90.55
##      16 comps 17 comps
## X      99.88   100.00
## Apps   93.42   93.89
```


```
pred.pcr <- predict(fit.pcr, test, ncomp=16) # min Cv at M=16
(err.pcr <- mean((test$Apps - pred.pcr)^2)) # test error
```

```
## [1] 1137877
```

Part f)

```
require(pls)
set.seed(1)
fit.pls <- plsr(Apps~., data=train, scale=TRUE, validation="CV")
validationplot(fit.pls, val.type="MSEP")
```

Apps



```
summary(fit.pls)
```

```
## Data:      X dimension: 388 17
## Y dimension: 388 1
## Fit method: kernelppls
## Number of components considered: 17
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV              4288    2217    2019    1761    1630    1533    1347
## adjCV           4288    2211    2012    1749    1605    1510    1331
##      7 comps  8 comps  9 comps 10 comps 11 comps 12 comps 13 comps
## CV          1309    1303    1286    1283    1283    1277    1271
## adjCV        1296    1289    1273    1270    1270    1264    1258
##      14 comps 15 comps 16 comps 17 comps
## CV           1270    1270    1270    1270
## adjCV        1258    1257    1257    1257
##
## TRAINING: % variance explained
##      1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps  8 comps
## X          27.21   50.73   63.06   65.52   70.20   74.20   78.62   80.81
## Apps       75.39   81.24   86.97   91.14   92.62   93.43   93.56   93.68
##      9 comps 10 comps 11 comps 12 comps 13 comps 14 comps 15 comps
## X          83.29   87.17   89.15   91.37   92.58   94.42   96.98
```


```
## Apps    93.76    93.79    93.83    93.86    93.88    93.89    93.89
##      16 comps  17 comps
## X        98.78   100.00
## Apps     93.89    93.89
```

```
pred.pls <- predict(fit.pls, test, ncomp=10) # min Cv at M=10
(err.pls <- mean((test$Apps - pred.pls)^2)) # test error
```

```
## [1] 1131661
```


Part g)

```
err.all <- c(err.lm, err.ridge, err.lasso, err.pcr, err.pls)
names(err.all) <- c("lm", "ridge", "lasso", "pcr", "pls")
barplot(err.all )
```



The test errors aren't much different. The ridge and lasso seem to perform slightly better while the PCR/PLS don't show any improvement from the full linear regression model.

```
plot(test$Apps, pred.lm)
```

EXERCISE 10:

Part a)

```
set.seed(1)
eps <- rnorm(1000)
xmat <- matrix(rnorm(1000*20), ncol=20)
betas <- sample(-5:5, 20, replace=TRUE)
betas[c(3,6,7,10,13,17)] <- 0
betas
```

```
## [1] -3 5 0 -5 2 0 0 4 5 0 4 1 0 5 -3 4 0 4 -2 1
```

```
y <- xmat %*% betas + eps
```

Part b)

```
set.seed(1)
trainid <- sample(1:1000, 100, replace=FALSE)
xmat.train <- xmat[trainid,]
xmat.test <- xmat[-trainid,]
y.train <- y[trainid,]
```

```

y.test <- y[-trainid,]
train <- data.frame(y=y.train, xmat.train)
test <- data.frame(y=y.test, xmat.test)

```

Part c)

```


require(leaps)

# predict function from chapter 6 labs
predict.regsubsets <- function(object, newdata, id, ...){
  form <- as.formula(object$call[[2]])
  mat <- model.matrix(form, newdata)
  coefi <- coef(object, id=id)
  xvars <- names(coefi)
  mat[,xvars] %*% coefi
}

regfit.full <- regsubsets(y~., data=train, nvmax=20)
err.full <- rep(NA, 20)
for(i in 1:20) {
  pred.full <- predict(regfit.full, train, id=i)
  err.full[i] <- mean((train$y - pred.full)^2)
}
plot(1:20, err.full, type="b", main="Training MSE", xlab="Number of Predictors")

```

Training MSE

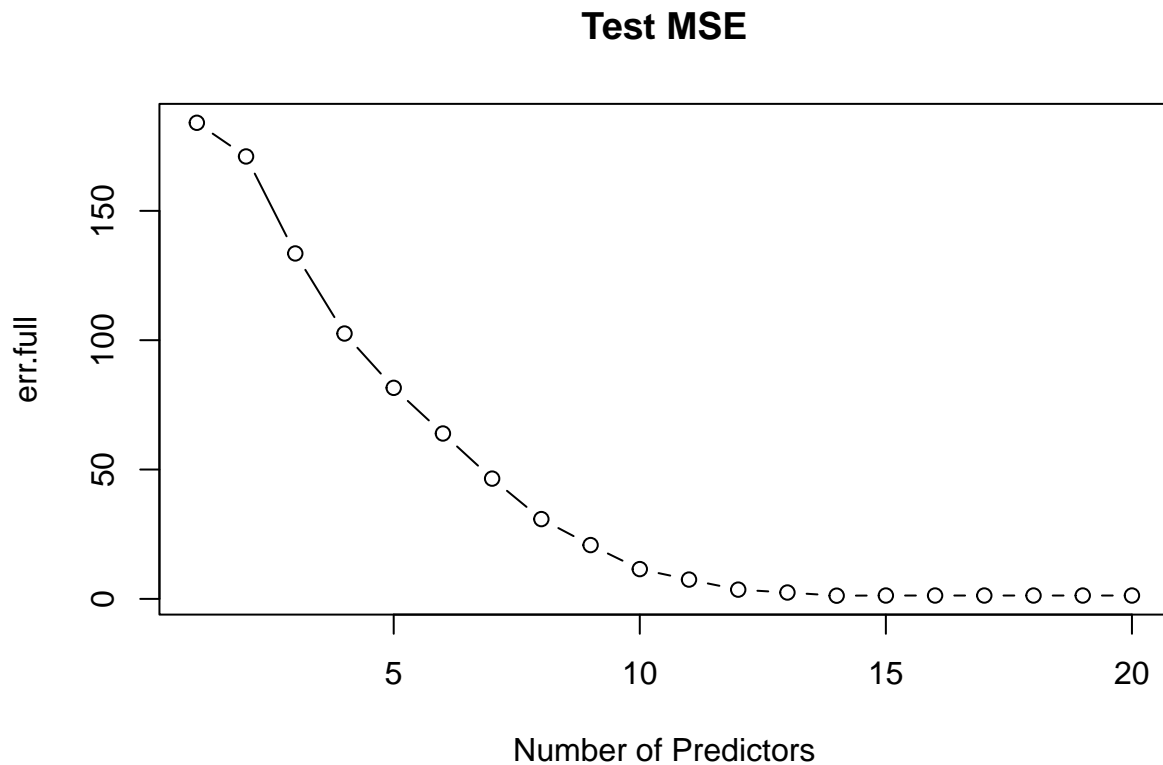


```
which.min(err.full) # min for train error should be at max pred count
```

```
## [1] 20
```

Part d)

```
err.full <- rep(NA, 20)
for(i in 1:20) {
  pred.full <- predict(regfit.full, test, id=i)
  err.full[i] <- mean((test$y - pred.full)^2)
}
plot(1:20, err.full, type="b", main="Test MSE", xlab="Number of Predictors")
```



Part e)

```
which.min(err.full) # optimal number of predictors from best subset
```

```
## [1] 14
```

Part f)

```
(coef.best <- coef(regfit.full, id=which.min(err.full)))
```

```
## (Intercept)      X1      X2      X4      X5      X8
## 0.002830717 -2.971273645 5.016860247 -5.032930569 2.011190672 4.043424774
##      X9      X11      X12      X14      X15      X16
## 5.166187040 4.002964100 1.038961952 4.950846519 -2.815120305 3.816000415
##      X18      X19      X20
## 4.146759315 -1.943471497 1.199621119
```

```
betas[betas != 0]
```

```
## [1] -3 5 -5 2 4 5 4 1 5 -3 4 4 -2 1
```

```
names(betas) <- paste0("X", 1:20)
merge(data.frame(beta=names(betas),betas), data.frame(beta=names(coef.best),coef.best), all.x=T, sort=F)
```


```
##      beta betas coef.best
## 1      X1      -3 -2.971274
## 2      X2       5  5.016860
## 3      X4      -5 -5.032931
## 4      X5       2  2.011191
## 5      X8       4  4.043425
## 6      X9       5  5.166187
## 7     X11       4  4.002964
## 8     X12       1  1.038962
## 9     X14       5  4.950847
## 10    X15      -3 -2.815120
## 11    X16       4  3.816000
## 12    X18       4  4.146759
## 13    X19      -2 -1.943471
## 14    X20       1  1.199621
## 15    X13       0      NA
## 16     X6       0      NA
## 17     X3       0      NA
## 18    X17       0      NA
## 19    X10       0      NA
## 20     X7       0      NA
```

The best subset model selected all the correct predictors

Part g)

```
err.coef <- rep(NA, 20)
for(i in 1:20) {
  coef.i <- coef(regfit.full, id=i)
  df.err <- merge(data.frame(beta=names(betas),betas), data.frame(beta=names(coef.i),coef.i), all.x=T)
  df.err[is.na(df.err[,3]),3] <- 0
  err.coef[i] <- sqrt(sum((df.err[,2] - df.err[,3])^2))
}
plot(1:20, err.coef, type="b", main="Coefficient Error", xlab="Number of Predictors")
points(which.min(err.coef), err.coef[which.min(err.coef)], col="red", pch=16)
```

Coefficient Error



The coefficient error plot shows a very similar plot to the test error plot

EXERCISE 11:

Part a)

```
require(leaps) # forward and backward selection
require(glmnet) # ridge and lasso
require(MASS) # Boston data set
data(Boston)

# split data into training and test sets
set.seed(1)
trainid <- sample(1:nrow(Boston), nrow(Boston)/2)
train <- Boston[trainid,]
test <- Boston[-trainid,]
xmat.train <- model.matrix(crim~., data=train)[,-1]
xmat.test <- model.matrix(crim~., data=test)[,-1]
str(Boston)
```

```
## 'data.frame':  506 obs. of  14 variables:
## $ crim : num  0.00632 0.02731 0.02729 0.03237 0.06905 ...
## $ zn : num  18 0 0 0 0 0 12.5 12.5 12.5 12.5 ...
## $ indus : num  2.31 7.07 7.07 2.18 2.18 2.18 7.87 7.87 7.87 7.87 ...
```

```
## $ chas : int 0 0 0 0 0 0 0 0 0 0 ...
## $ nox : num 0.538 0.469 0.469 0.458 0.458 0.458 0.524 0.524 0.524 0.524 ...
## $ rm : num 6.58 6.42 7.18 7 7.15 ...
## $ age : num 65.2 78.9 61.1 45.8 54.2 58.7 66.6 96.1 100 85.9 ...
## $ dis : num 4.09 4.97 4.97 6.06 6.06 ...
## $ rad : int 1 2 2 3 3 3 5 5 5 5 ...
## $ tax : num 296 242 242 222 222 222 311 311 311 311 ...
## $ ptratio: num 15.3 17.8 17.8 18.7 18.7 18.7 15.2 15.2 15.2 15.2 ...
## $ black : num 397 397 393 395 397 ...
## $ lstat : num 4.98 9.14 4.03 2.94 5.33 ...
## $ medv : num 24 21.6 34.7 33.4 36.2 28.7 22.9 27.1 16.5 18.9 ...
```

```
# ridge regression model
fit.ridge <- cv.glmnet(xmat.train, train$crim, alpha=0)
(lambda <- fit.ridge$lambda.min) # optimal lambda
```

```
## [1] 0.5919159
```

```
pred.ridge <- predict(fit.ridge, s=lambda, newx=xmat.test)
(err.ridge <- mean((test$crim - pred.ridge)^2)) # test error
```

```
## [1] 40.92777
```

```
predict(fit.ridge, s=lambda, type="coefficients")
```

```
## 14 x 1 sparse Matrix of class "dgCMatrix"
##          s1
## (Intercept) 14.702068319
## zn          0.035283661
## indus       -0.119976460
## chas        -0.616052143
## nox         -5.629356997
## rm          0.228001209
## age         -0.004314219
## dis         -0.768474303
## rad         0.434236779
## tax         0.003139323
## ptratio     -0.298647697
## black       -0.013823430
## lstat       0.262179351
## medv        -0.146028474
```

```
# lasso regression model
fit.lasso <- cv.glmnet(xmat.train, train$crim, alpha=1)
(lambda <- fit.lasso$lambda.min) # optimal lambda
```

```
## [1] 0.06805595
```

```
pred.lasso <- predict(fit.lasso, s=lambda, newx=xmat.test)
(err.lasso <- mean((test$crim - pred.lasso)^2)) # test error
```

```
## [1] 40.90173
```

```
predict(fit.lasso, s=lambda, type="coefficients")
```

```
## 14 x 1 sparse Matrix of class "dgCMatrix"
```

```
##              s1
## (Intercept) 17.65005513
## zn          0.03516255
## indus       -0.11838293
## chas        -0.43135144
## nox         -7.19578180
## rm          0.04271112
## age         .
## dis        -0.76801501
## rad         0.52430211
## tax         .
## ptratio    -0.35072332
## black      -0.01307754
## lstat       0.25559458
## medv       -0.14805010
```

```
# predict function from chapter 6 labs
```

```
predict.regsubsets <- function(object, newdata, id, ...){
  form <- as.formula(object$call[[2]])
  mat <- model.matrix(form, newdata)
  coefi <- coef(object, id=id)
  xvars <- names(coefi)
  mat[,xvars]%*%coefi
}
```

```
# forward selection
```

```
fit.fwd <- regsubsets(crim~., data=train, nvmax=ncol(Boston)-1)
(fwd.summary <- summary(fit.fwd))
```

```
## Subset selection object
```

```
## Call: regsubsets.formula(crim ~ ., data = train, nvmax = ncol(Boston) -
##      1)
```

```
## 13 Variables (and intercept)
```

```
##      Forced in Forced out
```

```
## zn          FALSE      FALSE
```

```
## indus       FALSE      FALSE
```

```
## chas        FALSE      FALSE
```

```
## nox         FALSE      FALSE
```

```
## rm          FALSE      FALSE
```

```
## age         FALSE      FALSE
```

```
## dis         FALSE      FALSE
```

```
## rad         FALSE      FALSE
```

```
## tax         FALSE      FALSE
```

```
## ptratio     FALSE      FALSE
```

```
## black       FALSE      FALSE
```


```
## lstat       FALSE      FALSE
```

```
## medv        FALSE      FALSE
```

```
## 1 subsets of each size up to 13
```

```
## Selection Algorithm: exhaustive
##          zn  indus chas nox rm  age dis rad tax ptratio black lstat medv
## 1  ( 1 )  " " " "  " " " " " " " " " " " " " " " " " " " " " " " "
## 2  ( 1 )  " " " "  " " " " " " " " " " " " " " " " " " " " " " " "
## 3  ( 1 )  " " " "  " " " " " " " " " " " " " " " " " " " " " " " "
## 4  ( 1 )  "*" " "  " " " " " " " " " " " " " " " " " " " " " " " "
## 5  ( 1 )  " " "*"  " " " " " " " " " " " " " " " " " " " " " " " "
## 6  ( 1 )  "*" "*"  " " " " " " " " " " " " " " " " " " " " " " " "
## 7  ( 1 )  "*" "*"  " " " " " " " " " " " " " " " " " " " " " " " "
## 8  ( 1 )  "*" " "  " " " " " " " " " " " " " " " " " " " " " " " "
## 9  ( 1 )  "*" "*"  " " " " " " " " " " " " " " " " " " " " " " " "
## 10 ( 1 )  "*" "*"  " " " " " " " " " " " " " " " " " " " " " " " "
## 11 ( 1 )  "*" "*"  "*" " " " " " " " " " " " " " " " " " " " " " "
## 12 ( 1 )  "*" "*"  "*" " " " " " " " " " " " " " " " " " " " " " "
## 13 ( 1 )  "*" "*"  "*" " " " " " " " " " " " " " " " " " " " " " "
```

```
err.fwd <- rep(NA, ncol(Boston)-1)
for(i in 1:(ncol(Boston)-1)) {
  pred.fwd <- predict(fit.fwd, test, id=i)
  err.fwd[i] <- mean((test$crim - pred.fwd)^2)
}
plot(err.fwd, type="b", main="Test MSE for Forward Selection", xlab="Number of Predictors")
```



```
which.min(err.fwd)
```

```
## [1] 1
```



```


# backward selection
fit.bwd <- regsubsets(crim~., data=train, nvmax=ncol(Boston)-1)
(bwd.summary <- summary(fit.bwd))

## Subset selection object
## Call: regsubsets.formula(crim ~ ., data = train, nvmax = ncol(Boston) -
##      1)
## 13 Variables (and intercept)
##      Forced in Forced out
## zn          FALSE      FALSE
## indus        FALSE      FALSE
## chas          FALSE      FALSE
## nox           FALSE      FALSE
## rm            FALSE      FALSE
## age           FALSE      FALSE
## dis           FALSE      FALSE
## rad           FALSE      FALSE
## tax           FALSE      FALSE
## ptratio       FALSE      FALSE
## black         FALSE      FALSE
## lstat         FALSE      FALSE
## medv          FALSE      FALSE
## 1 subsets of each size up to 13
## Selection Algorithm: exhaustive
##      zn indus chas nox rm age dis rad tax ptratio black lstat medv
## 1 ( 1 ) " " " " " " " " " " " " " " " " " " " " " "
## 2 ( 1 ) " " " " " " " " " " " " " " " " " " " " "
## 3 ( 1 ) " " " " " " " " " " " " " " " " " " " " "
## 4 ( 1 ) "*" " " " " " " " " " " " " " " " " " " " "
## 5 ( 1 ) " " "*" " " " " " " " " " " " " " " " " " "
## 6 ( 1 ) "*" "*" " " " " " " " " " " " " " " " " " "
## 7 ( 1 ) "*" "*" " " " " " " " " " " " " " " " " " "
## 8 ( 1 ) "*" " " " " " " " " " " " " " " " " " " " "
## 9 ( 1 ) "*" "*" " " " " " " " " " " " " " " " " " "
## 10 ( 1 ) "*" "*" " " " " " " " " " " " " " " " " " "
## 11 ( 1 ) "*" "*" "*" " " " " " " " " " " " " " " " "
## 12 ( 1 ) "*" "*" "*" " " " " " " " " " " " " " " " "
## 13 ( 1 ) "*" "*" "*" " " " " " " " " " " " " " " " "

err.bwd <- rep(NA, ncol(Boston)-1)
for(i in 1:(ncol(Boston)-1)) {
  pred.bwd <- predict(fit.bwd, test, id=i)
  err.bwd[i] <- mean((test$crim - pred.bwd)^2)
}
plot(err.bwd, type="b", main="Test MSE for Backward Selection", xlab="Number of Predictors")

```

Test MSE for Backward Selection



```
which.min(err.bwd)
```

```
## [1] 1
```

```
par(mfrow=c(3,2))
```

```
min.cp <- which.min(fwd.summary$cp)
```

```
plot(fwd.summary$cp, xlab="Number of Poly(X)", ylab="Forward Selection Cp", type="l")
points(min.cp, fwd.summary$cp[min.cp], col="red", pch=4, lwd=5)
```

```
min.cp <- which.min(bwd.summary$cp)
```

```
plot(bwd.summary$cp, xlab="Number of Poly(X)", ylab="Backward Selection Cp", type="l")
points(min.cp, bwd.summary$cp[min.cp], col="red", pch=4, lwd=5)
```

```
min.bic <- which.min(fwd.summary$bic)
```

```
plot(fwd.summary$bic, xlab="Number of Poly(X)", ylab="Forward Selection BIC", type="l")
points(min.bic, fwd.summary$bic[min.bic], col="red", pch=4, lwd=5)
```

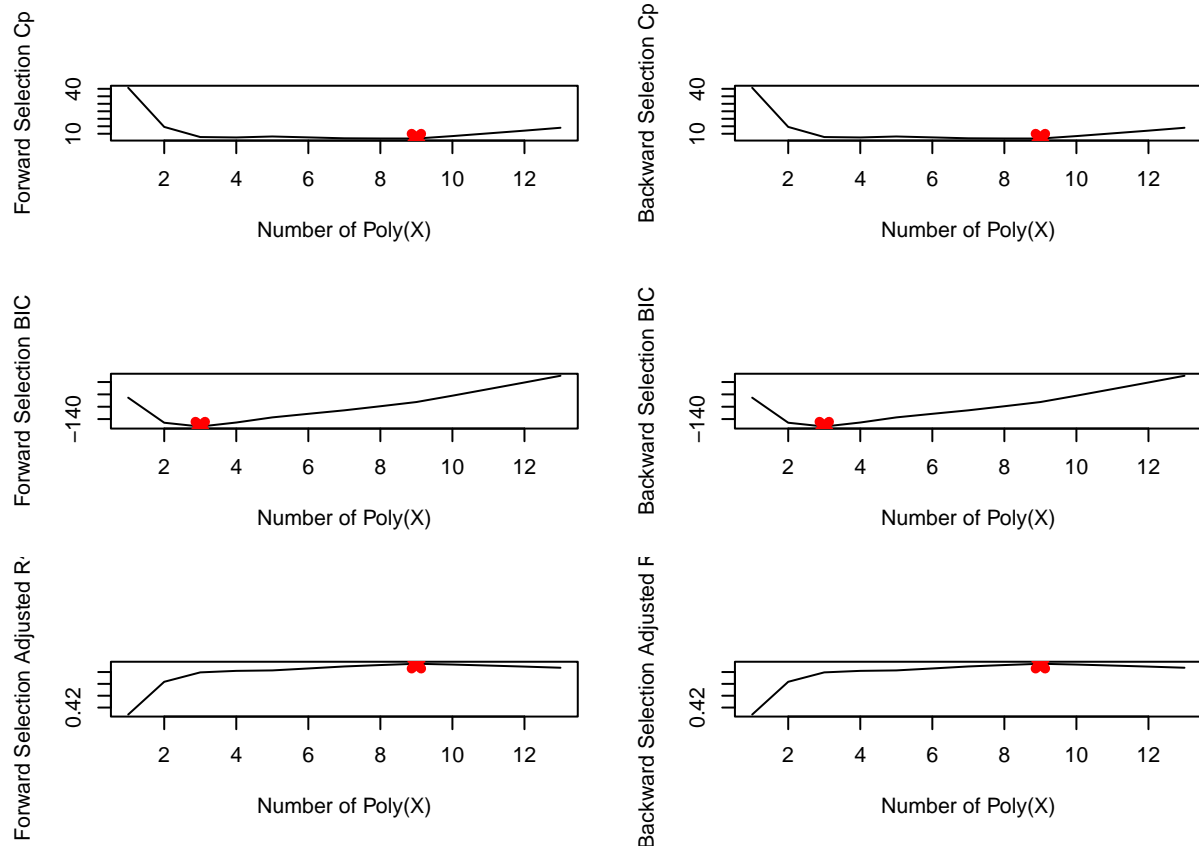
```
min.bic <- which.min(bwd.summary$bic)
```

```
plot(bwd.summary$bic, xlab="Number of Poly(X)", ylab="Backward Selection BIC", type="l")
points(min.bic, bwd.summary$bic[min.bic], col="red", pch=4, lwd=5)
```

```
min.adj2 <- which.max(fwd.summary$adj2)
```

```
plot(fwd.summary$adj2, xlab="Number of Poly(X)", ylab="Forward Selection Adjusted R^2", type="l")
points(min.adj2, fwd.summary$adj2[min.adj2], col="red", pch=4, lwd=5)
```

```
min.adj2 <- which.max(bwd.summary$adjr2)
plot(bwd.summary$adjr2, xlab="Number of Poly(X)", ylab="Backward Selection Adjusted R^2", type="l")
points(min.adj2, bwd.summary$adjr2[min.adj2], col="red", pch=4, lwd=5)
```



Part b)

```
err.ridge
```

```
## [1] 40.92777
```

```
err.lasso
```

```
## [1] 40.90173
```

```
err.fwd
```

```
## [1] 40.14557 41.87706 42.40901 42.45745 43.03836 42.13258 41.25016 41.28957
## [9] 41.49271 41.50577 41.48839 41.46692 41.54639
```

```
err.bwd
```

```
## [1] 40.14557 41.87706 42.40901 42.45745 43.03836 42.13258 41.25016 41.28957
## [9] 41.49271 41.50577 41.48839 41.46692 41.54639
```

Probably choose the lasso model because its test MSE is close to best and eliminates some predictors to reduce model complexity

Part c)

No because not all the predictors add much value to the model